andrewbelyaev2164@gmail.com
github.com/andreibeliaev

Andrei Beliaev

andreibeliaev.github.io

linkedin.com/in/andreibeliaev
+1 (217) 419 9945

Experience

Epivara, Inc., BioTech — Champaign, IL
AI/ML Software Engineer — *May 2025 – Present*

- Built and pretrained DINOv3 ViT + CNN-adapter model capturing tissue-scale structure down to sub-nuclear detail: +18.5% mSA and +4.6% F1 over a state-of-the-art baseline on Leydig-cell histology; underpinned a toxicology study.
- Engineered a 16× larger field of view at 2× lower inference cost and 3.7× faster training over a state-of-the-art baseline, supplying the anatomical context needed to reject biologically implausible classifications across ∼3/4 of surveyed tissue endpoints.
- Built memory-based animal segmentation and tracking with Cutie, custom state management, and cross-camera checks; fine-tuned DINOv3 for re-identification (92.3% macro F1). Trained V-JEPA 2 heads for video-based behavior classification.
- Built a multi-camera 3D pose pipeline for animal behavior analysis using partial 2D/3D labels and temporal refinement; achieved 5.3 mm joint error on held-out frames, versus 8.4 mm for Lightning Pose 3D and 16.4 mm for fine-tuned s-DANNCE.
- Shipped a research-data *analysis platform* ↗ with custom LLM orchestration and MCP tools for schema exploration, SQL, sandboxed Python, and reporting; built its FastAPI backend and React/TypeScript interface.
- Developed a human-in-the-loop ingestion agent that reduced experimental-data insertion from hours to minutes.

Tinkoff, FinTech — Moscow, Russia
Data Scientist, Compliance — *Feb 2022 – Oct 2022*

- Contributed to feature development for a new Anti-Money Laundering model. The model increased fully automated case resolution from 45% to 60% and cut annual costs by $1.7M.
- Developed SQL/Python/Spark pipelines and dashboards that improved fraud and compliance analysis, cutting 1,000 hours of manual work annually.

Data Analyst, Insurance — *Mar 2021 – Feb 2022*

- Engineered a real-time quality measurement framework that merged backend logs, frontend events, and database tables, exposing live SRE metrics in Grafana and raising product quality (aggregated metric) by 18% while saving $60K monthly.
- Delivered ML-ready datasets for churn prediction, ran A/B tests, and worked with product and engineering teams to diagnose production issues.

Open Source

HistoForge: AI Platform for Pathology [Link]

- Built an end-to-end platform (QuPath Java plugin + Modal GPU inference runtime + Hugging Face model hub + pip package) that brings foundation-model tissue analysis to pathologists and researchers.
- Pretrained a large field-of-view custom DINOv3 + CNN-adapter model on 16 public cell-segmentation datasets; built 7 nuclei- and structure-level modules trained on public data.
- Packaged training and inference so labs can fine-tune, run, and share their own models and pipelines, targeting reproducible and accessible quantitative pathology for oncology and toxicology research.

Projects

Improving System-2 Thinking in Energy-Based Transformers [Link]

- Ran MCMC sampling ablations in Energy-Based Transformers, comparing perplexity across pretraining configurations.

Image Immunization Against Diffusion-Based Editing

- Implemented image-immunization methods from CVPR 2024 and ICLR 2026, using timestep-universal gradients and learned perturbation networks for Stable Diffusion.

Entropy-aware sampling in vLLM [Link]

- Implemented entropy-aware token sampling in vLLM with GPU-batched lookahead to control diversity by penalizing entropy-reducing tokens in speculative decoding.

Education

University of Illinois Urbana–Champaign — Champaign, IL
Master of Computer Science | Graduate Research Assistant — May 2026

Bauman Moscow State Technical University — Moscow, Russia
Specialist in Autonomous Informational and Control Systems — Jun 2022

Skills

Languages: Python (FastAPI, SQLAlchemy, asyncio), SQL, TypeScript (React, Zustand, shadcn/ui), Java, C++, MATLAB
AI/ML: PyTorch, LangGraph, vLLM, Scikit-learn, W&B, tool and agent orchestration, HITL, RAG
Infra/Data: AWS (EC2, S3, IAM, ECS, VPC), Modal, Docker, PySpark, PostgreSQL, HPC (SLURM)